

## Sta314 HW2 Problem 2

2023-10-17

#Problem 2

```
library(MASS)
set.seed(0)
```

##Part 1 Generate the training data with  $\rho = 0.1$  and  $n = 100$

```
rho = 0.1
n = 100
beta = c(0.5,0.5)
mean = c(0,0,0)
sigma = matrix(c(1,0,0,0,1,rho,0,rho,1), nrow = 3)
X = mvrnorm(n, mean, sigma)
epsilon = rnorm(n) # since epsilon(i) with  $1 \leq i \leq n$  i.i.d. from  $N(0,1)$ 
y = X[c(1,2)]%*% beta + epsilon
sample = data.frame(cbind(X,y))
```

```
f_hat_1 = lm(y~V1 + V2, data = sample)
f_hat_2 = lm(y~V1 + V2 + V3, data = sample)
f_hat_3 = lm(y~V1 + V3, data = sample)
```

```
MSE_model_1 = mean((y - predict(f_hat_1, data = sample))**2)
MSE_model_2 = mean((y - predict(f_hat_2, data = sample))**2)
MSE_model_3 = mean((y - predict(f_hat_3, data = sample))**2)
```

```
cat("MSE Model 1:", MSE_model_1, "\n")
```

## MSE Model 1: 0.8737418

```
cat("MSE Model 2:", MSE_model_2, "\n")
```

## MSE Model 2: 0.860581

```
cat("MSE Model 3:", MSE_model_3, "\n")
```

## MSE Model 3: 0.8725934

Get the same conclusions as in Problem 1. Since the model 2 has 3 features, such that the value of  $\text{MSE}(\hat{f}_2)$  is the lowest.

##Part 2 Repeat part 1 above  $N = 100$

```
n = 100
mean = c(0, 0, 0)
sigma = matrix(c(1, 0.5, 0.5, 0.5, 1, 0.5, 0.5, 0.5, 1), nrow = 3)
beta = c(0.5, 0.5)
count = 0
```

```
for (i in 1:100) { #Since need to repeat 100 times, trying to use for i in range
  X = as.data.frame(mvrnorm(n, mean, sigma))
```

```

epsilon = rnorm(n)
y = X$V1 * beta[1] + X$V2 * beta[2] + epsilon
sample = data.frame(X, y)

f_hat_1 = lm(y ~ V1 + V2, data = sample)
f_hat_2 = lm(y ~ V1 + V2 + V3, data = sample)
f_hat_3 = lm(y ~ V1 + V3, data = sample)

MSE_model_1 = mean((y - predict(f_hat_1, data = sample))^2)
MSE_model_2 = mean((y - predict(f_hat_2, data = sample))^2)
MSE_model_3 = mean((y - predict(f_hat_3, data = sample))^2)

if (MSE_model_2 < MSE_model_1) {
  count = count + 1
}
}

cat("Count:", count, "\n")

```

## Count: 100

For part 2, we need to repeat part 1 100 times, trying to use for loop function, set up the count = 0 at the beginning, since part 1 we get the conclusion that the model 2 has 3 features, such that the value of  $MSE(f\_hat\_2)$  is the lowest, set up the if statement, which is if  $(MSE\_model\_2 < MSE\_model\_1)$ , finally, get the count = 100, which means  $MSE\_model\_2$  always lower than  $MSE\_model\_1$ , but we can not compare  $MSE\_model\_1$  and  $MSE\_model\_3$ , with the same conclusion in Question 1 part 3.

##Part 3 Repeat part 2 above rho = 0.95

```

rho = 0.95
n = 100
mean = c(0, 0, 0)
sigma = matrix(c(1, 0.5, 0.5, 0.5, 1, 0.5, 0.5, 0.5, 1), nrow = 3)
beta = c(0.5, 0.5)
count1v2 = count1v3 = 0

for (i in 1:100) { #Since need to repeat 100 times, trying to use for i in range
  X = as.data.frame(mvrnorm(n, mean, sigma))
  epsilon = rnorm(n)
  y = X$V1 * beta[1] + X$V2 * beta[2] + epsilon
  sample = data.frame(X, y)

  f_hat_1 = lm(y ~ V1 + V2, data = sample)
  f_hat_2 = lm(y ~ V1 + V2 + V3, data = sample)
  f_hat_3 = lm(y ~ V1 + V3, data = sample)

  MSE_model_1 = mean((y - predict(f_hat_1, data = sample))^2)
  MSE_model_2 = mean((y - predict(f_hat_2, data = sample))^2)
  MSE_model_3 = mean((y - predict(f_hat_3, data = sample))^2)

  if (MSE_model_2 < MSE_model_1) {
    count1v2 = count1v2 + 1
  }

  if (MSE_model_1 < MSE_model_3) {

```

```

        count1v3 = count1v3 + 1
    }
}
print(paste("MSE Model 2 was lower than Model 1 in", count1v2,
           "out of 100 trials.))

## [1] "MSE Model 2 was lower than Model 1 in 100 out of 100 trials."

print(paste("MSE Model 1 was lower than Model 3 in", count1v3,
           "out of 100 trials.))

## [1] "MSE Model 1 was lower than Model 3 in 99 out of 100 trials."

```

Since  $\rho$  represent the correlation between  $X_2$  and  $X_3$ . From part 1  $\rho = 0.1$ , and part 3  $\rho = 0.95$ , the result shows that MSE Model 2 was lower than Model 1 in 100 out of 100 trials, which same result with different  $\rho$ .  $\rho = 0.95$ , which means  $X_2$  and  $X_3$  both have high correlation, which means  $\hat{f}_1$  and  $\hat{f}_3$  essential is the same, such that the result shows that MSE Model 1 was lower than Model 3 in 99 out of 100 trials

# Sta314 HW2 Problem 3

2023-10-19

##1 Generate a data set

```
n = 1000
p = 20
beta = c(rep(2,3), rep(0.5,2), rep(0,15))
```

```
X = matrix(rnorm(n*p), nrow = n)
epsilon = rnorm(n)
y = X %*% beta + epsilon
```

*# Create data frame*

```
data = data.frame(y, X)
```

##2 Randomly split dataset

```
set.seed(0)
```

```
index = sample(1:n, 100)
```

```
train.dataset = data[index, ] #training set containing 100 observations
```

```
test.dataset = data[-index, ] #test set containing 900 observations
```


##3 Perform best subset selection

```
library(leaps)
```

```
regfit.best = regsubsets(y ~ . , data = train.dataset, nvmax = 20)
```

```
reg_summary = summary(regfit.best)
```

```
plot(reg_summary$rss/length(index), xlab = 'number of predictors',
     ylab = 'Training MSE', type = 'l')
```



##4 Plot the test set MSE associated with the best model of each size


```
Test_MSE = rep(NA, 20)

for (i in 1:p) {
  coef_i = coef(regfit.best, id = i)
  has_intercept <- "(Intercept)" %in% names(coef_i)

  if (has_intercept) {
    predictors <- setdiff(names(coef_i), "(Intercept)")
  } else {
    predictors <- names(coef_i)
  }

  X_test = cbind(1, as.matrix(test.dataset[, predictors]))
  pred_i = X_test %*% coef_i
  Test_MSE[i] = mean((test.dataset$y - pred_i)^2)
}

plot(Test_MSE, type = "l", xlab = "Number of Predictors", ylab = "Test MSE")
```



##5 For which model size do the training set MSE and test set MSE

```
Train_MSE = reg_summary$rss
which.min(Train_MSE)
```

```
## [1] 20
```

```
which.min(Test_MSE)
```

```
## [1] 5
```

The model size is 20, the training MSE reaches its lowest. When the model size is 5, the test MSE is at its minimum. Model with more predictors fit the training data more closely. Such that selecting a model based on the training MSE can result in overfitting, but if using the test MSE as a guide often leads to a model that's more representative of the true model.

##6 How does the model at which the test set MSE is minimized

```
coef(regfit.best, id = 5 )
```

```
## (Intercept)          X1          X2          X3          X4          X5
##  0.06068251  2.11667329  2.19704836  2.01389150  0.34128119  0.51954004
```

Since in the part 1, we set the true model for the beta,  $\beta_1 - \beta_3 = 2$ , and  $\beta_4 - \beta_5 = 0.5$ ,  $\beta_6 - \beta_{20} = 0$ . The result shows that coefficient of  $X1\_hat - X3\_hat$  more close to 2, and coefficient  $X4\_hat - X5\_hat$  close to 0.5. Such that selecting a model based on the test MSE will more close to true model. Since the epsilon exist, which the test set not exactly 100% like true model.

##7 Create a plot displaying

```
result = rep(NA, p)
```

```
for(i in 1:p) {
  beta.hati = rep(0, p + 1)
  names(beta.hati) = c("(Intercept)", colnames(test.dataset)[-1])

  coef_i = coef(regfit.best, id = i)
```


```

beta.hati[names(coef_i)] = coef_i

result[i] <- sqrt(sum((beta - beta.hati[-1])^2))
}

plot(result, xlab = "Number of predictors", ylab = "l2 norm", type = "l")

```




Since the result shows that at size 5, has the minimum value. Which is similar with part 4 test MSE plot.

##8 Repeat parts 3-6 for forward

```

#3
library(leaps)
regfit.forward = regsubsets(y ~ ., data = train.dataset, nvmax = 20,
                           method = 'forward')
reg_summary = summary(regfit.forward)
plot(reg_summary$rss/length(index), xlab = 'number of predictors',
     ylab = 'Training MSE', type = 'l')

```



```
#4
Test_MSE = rep(NA, 20)


for (i in 1:p) {
  coef_i = coef(regfit.forward, id = i)
  has_intercept <- "(Intercept)" %in% names(coef_i)

  if (has_intercept) {
    predictors <- setdiff(names(coef_i), "(Intercept)")
  } else {
    predictors <- names(coef_i)
  }

  X_test = cbind(1, as.matrix(test.dataset[, predictors]))
  pred_i = X_test %*% coef_i
  Test_MSE[i] = mean((test.dataset$y - pred_i)^2)
}

plot(Test_MSE, type = "l", xlab = "Number of Predictors", ylab = "Test MSE")
```





```
#5
Train_MSE = reg_summary$rss
which.min(Train_MSE)
```

```
## [1] 20
```

```
which.min(Test_MSE)
```

```
## [1] 5
```


The model size is 20, the training MSE reaches its lowest. When the model size is 5, the test MSE is at its minimum.

```
#6
coef(regfit.forward, id =5 )
```

```
## (Intercept)          X1          X2          X3          X4          X5
##  0.06068251  2.11667329  2.19704836  2.01389150  0.34128119  0.51954004
```

```
##9 Repeat parts 3-6 for backward
```

```
#3
library(leaps)
regfit.backward = regsubsets(y ~ . , data = train.dataset, nvmax = 20,
                             method = 'backward')
reg_summary = summary(regfit.backward)
plot(reg_summary$rss/length(index), xlab = 'number of predictors',
     ylab = 'Training MSE', type = 'l')
```




```
#4
Test_MSE = rep(NA, 20)

for (i in 1:p) {
  coef_i = coef(regfit.backward, id = i)
  has_intercept <- "(Intercept)" %in% names(coef_i)

  if (has_intercept) {
    predictors <- setdiff(names(coef_i), "(Intercept)")
  } else {
    predictors <- names(coef_i)
  }

  X_test = cbind(1, as.matrix(test.dataset[, predictors]))
  pred_i = X_test %*% coef_i
  Test_MSE[i] = mean((test.dataset$y - pred_i)^2)
}

plot(Test_MSE, type = "l", xlab = "Number of Predictors", ylab = "Test MSE")
```



```
#5
Train_MSE = reg_summary$RSS
which.min(Train_MSE)
```

```
## [1] 20
```

```
which.min(Test_MSE)
```

```
## [1] 5
```

The model size is 20, the training MSE reaches its lowest. When the model size is 5, the test MSE is at its minimum.

```
#6
coef(regfit.backward, id = 5)
```

```
## (Intercept)      X1      X2      X3      X4      X5
## 0.06068251 2.11667329 2.19704836 2.01389150 0.34128119 0.51954004
```

# Sta314 Hw3 Problem 4

2023-10-18

a~d

```
p = 50
n = 1100
set.seed(0)
X = matrix(rnorm(n * p), nrow = n,)
epsilon = rnorm(n)
beta = c(rep(2, 5), rep(0, 45)) # since p = 50
y = X %*% beta + epsilon

index = sample(1:n, 100) # randomly split the data: training 100, test 1000

train_x = X[index, ]
train_y = y[index]

test_x = X[-index, ]
test_y = y[-index]

grid = 10^seq(10, -2, length = 100)
```

## 1 Fit both the ridge regression and the lasso

```
library(glmnet)

## Loading required package: Matrix
## Loaded glmnet 4.1-8

# Ridge
ridge_cv = cv.glmnet(train_x, train_y, alpha = 0, lambda = grid)
best_lam = ridge_cv$lambda.min
ridge.mod = glmnet(train_x, train_y, alpha = 0, lambda = best_lam)
pred.ridge = predict(ridge.mod, test_x, lambda = best_lam)
MSE.ridge = mean((test_y - pred.ridge)**2)
MSE.ridge

## [1] 1.87885
```

lasso

```
# Lasso
lasso_cv = cv.glmnet(train_x, train_y, alpha = 1, lambda = grid)
best_lam = lasso_cv$lambda.min
lasso.mod = glmnet(train_x, train_y, alpha = 1, lambda = best_lam)
pred.lasso = predict(lasso.mod, test_x, lambda = best_lam)
```

```
MSE.lasso = mean((test_y - pred.lasso)^2)
MSE.lasso
```

```
## [1] 1.150411
```

Lasso leads to smaller test set MSE.

## 2 Repeat steps a~d for generating the data by using different seeds

```
Cal_MSE <- function(seed, p = 50, n = 1100) {
  set.seed(seed)
  X = matrix(rnorm(n * p), nrow = n)
  epsilon = rnorm(n)
  beta = c(rep(2, 5), rep(0, 45)) # since p = 50
  y = X %*% beta + epsilon

  index = sample(1:n, 100)
  train_x = X[index, ]
  train_y = y[index]
  test_x = X[-index, ]
  test_y = y[-index]
  grid = 10^seq(10, -2, length = 100)

  # Ridge
  ridge_cv = cv.glmnet(train_x, train_y, alpha = 0, lambda = grid)
  best_lam = ridge_cv$lambda.min
  ridge.mod = glmnet(train_x, train_y, alpha = 0, lambda = best_lam)
  pred.ridge = predict(ridge.mod, test_x, lambda = best_lam)
  MSE.ridge = mean((test_y - pred.ridge)**2)
  MSE.ridge


  # Lasso
  lasso_cv = cv.glmnet(train_x, train_y, alpha = 1, lambda = grid)
  best_lam = lasso_cv$lambda.min
  lasso.mod = glmnet(train_x, train_y, alpha = 1, lambda = best_lam)
  pred.lasso = predict(lasso.mod, test_x, lambda = best_lam)
  MSE.lasso = mean((test_y - pred.lasso)^2)
  MSE.lasso

  return(c(MSE_ridge = MSE.ridge , MSE_lasso = MSE.lasso ))
}

MSE_ridge_list = c(MSE.ridge)
MSE_lasso_list = c(MSE.lasso)

for (i in 2:50){ #repeat ,since the first MSE for ridge and lasso exist
  MSE = Cal_MSE(i)
  MSE_ridge_list = c(MSE_ridge_list, MSE[1])
  MSE_lasso_list = c(MSE_lasso_list, MSE[2])
}

boxplot(MSE_ridge_list, MSE_lasso_list, names= c('Ridge', 'Lasso'),
        ylab = 'Test MSE',
        col='green',
        border ='blue')
```



Using the boxplot, Ridge's MSE large than Lasso's MSE, also Ridge's MSE variance larger than Lasso's MSE.

### 3 Redo part 1 and 2 by using $\beta_j = 0.5$ for all $j=1, \dots, 50$ in step b

```
Cal_MSE <- function(seed, p = 50, n = 1100) {
  set.seed(seed)
  X = matrix(rnorm(n * p), nrow = n)
  epsilon = rnorm(n)
  beta = c(rep(0.5, 50)) # since p = 50
  y = X %*% beta + epsilon

  index = sample(1:n, 100)
  train_x = X[index, ]
  train_y = y[index]
  test_x = X[-index, ]
  test_y = y[-index]
  grid = 10^seq(10, -2, length = 100)

  # Ridge
  ridge_cv = cv.glmnet(train_x, train_y, alpha = 0, lambda = grid)
  best_lam = ridge_cv$lambda.min
  ridge.mod = glmnet(train_x, train_y, alpha = 0, lambda = best_lam)
  pred.ridge = predict(ridge.mod, test_x, lambda = best_lam)
  MSE.ridge = mean((test_y - pred.ridge)**2)
  MSE.ridge

  # Lasso
  lasso_cv = cv.glmnet(train_x, train_y, alpha = 1, lambda = grid)
  best_lam = lasso_cv$lambda.min
  lasso.mod = glmnet(train_x, train_y, alpha = 1, lambda = best_lam)
  pred.lasso = predict(lasso.mod, test_x, lambda = best_lam)
  MSE.lasso = mean((test_y - pred.lasso)**2)
  MSE.lasso

  return(c(MSE_ridge = MSE.ridge , MSE_lasso = MSE.lasso ))
}
```

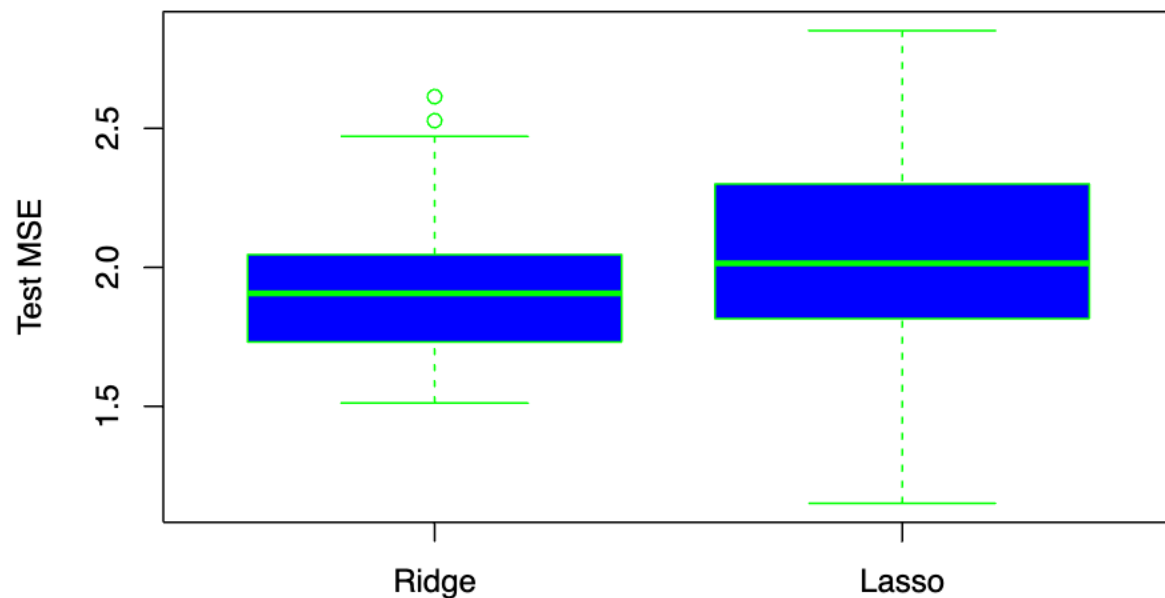
```

MSE_ridge_list = c(MSE.ridge)
MSE_lasso_list = c(MSE.lasso)

for (i in 2:50){ #repeat ,since the first MSE for ridge and lasso exist
  MSE = Cal_MSE(i)
  MSE_ridge_list = c(MSE_ridge_list, MSE[1])
  MSE_lasso_list = c(MSE_lasso_list, MSE[2])
}

boxplot(MSE_ridge_list, MSE_lasso_list, names= c('Ridge', 'Lasso'),
        ylab = 'Test MSE',
        col='blue',
        border = 'green')

```



Since change the beta to 0.5, Lasso's MSE bigger than Ridge's MSE. The variance of Lasso MSE slightly large than Ridge's MSE.